

Nextflow with Singularity Containers



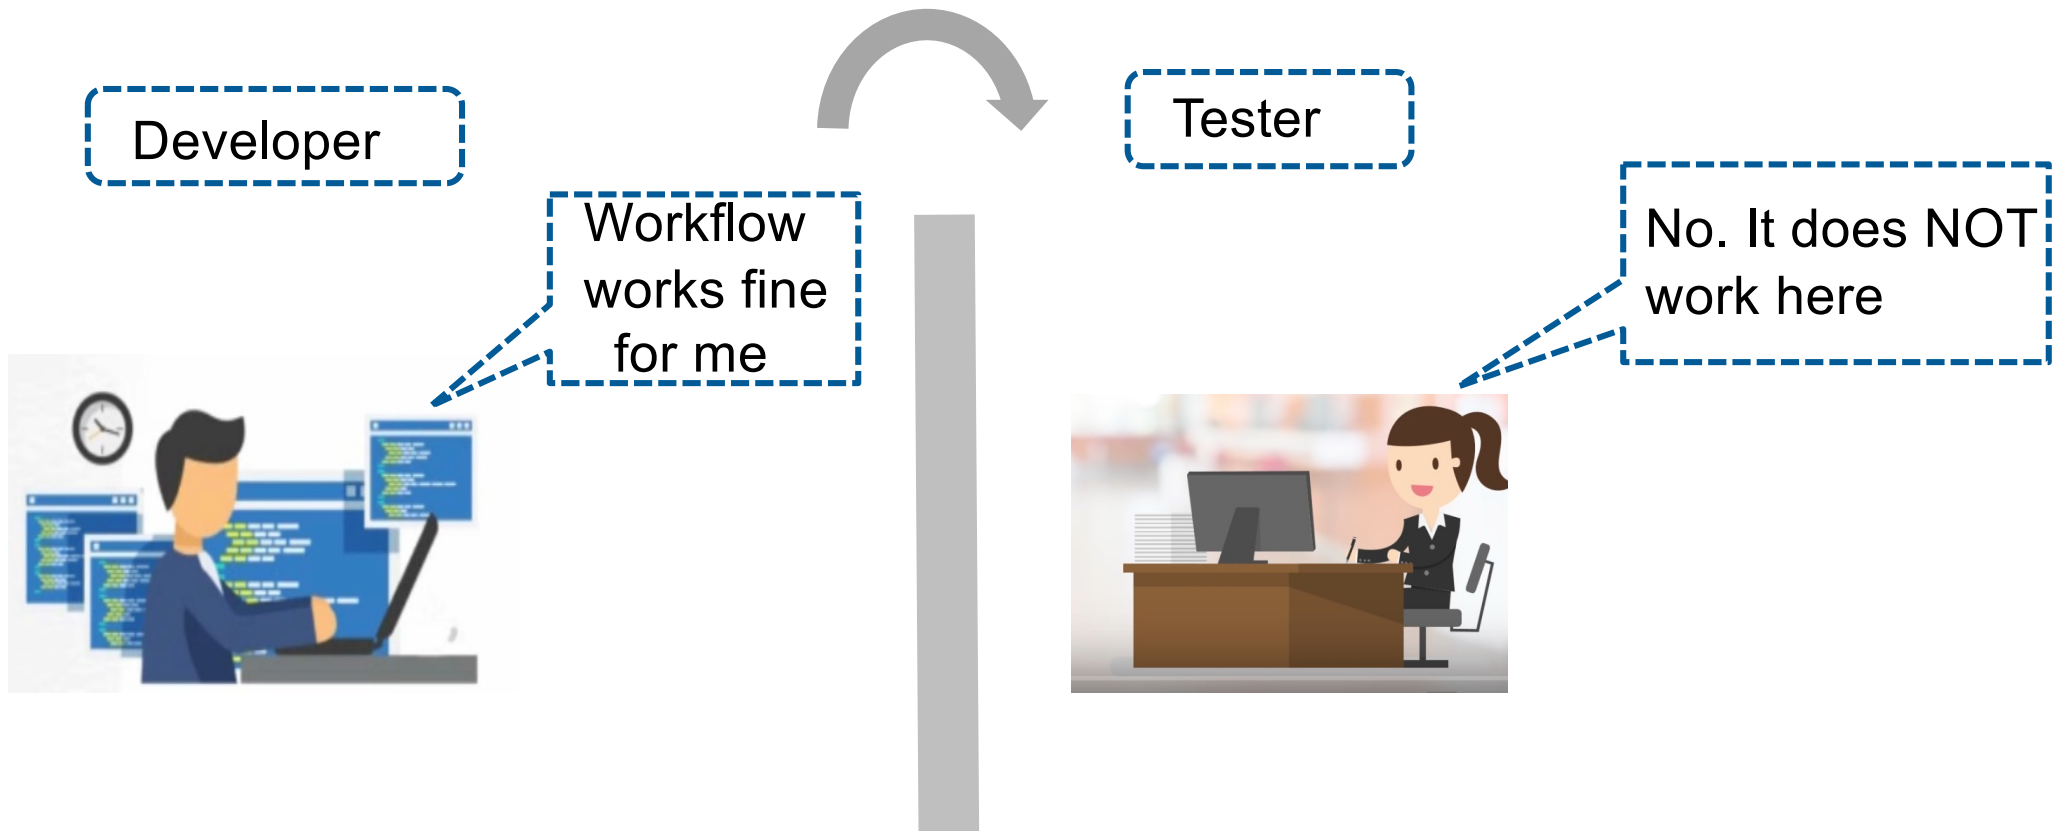
CSC – Suomalainen tutkimuksen, koulutuksen, kulttuurin ja julkishallinnon ICT-osaamiskeskus

Outline

- Why containers in workflows?
- Configuring singularity with Nextflow
- Running Nextflow workflows on Puhti
- Reporting and visualisation of workflows
- Working with *nf-core* workflows

Why Containers in Workflows?

Why Containers in Workflows?



Nextflow workflows with Containers



- Built-in integration with containers
- Advantages
 - Maintainability
 - Portability
 - Reproducibility
- Popular container choices
 - Docker
 - Singularity
 - (conda)

Which software to choose for Nextflow ?

nextflow

+



Docker

OR



Conda

OR

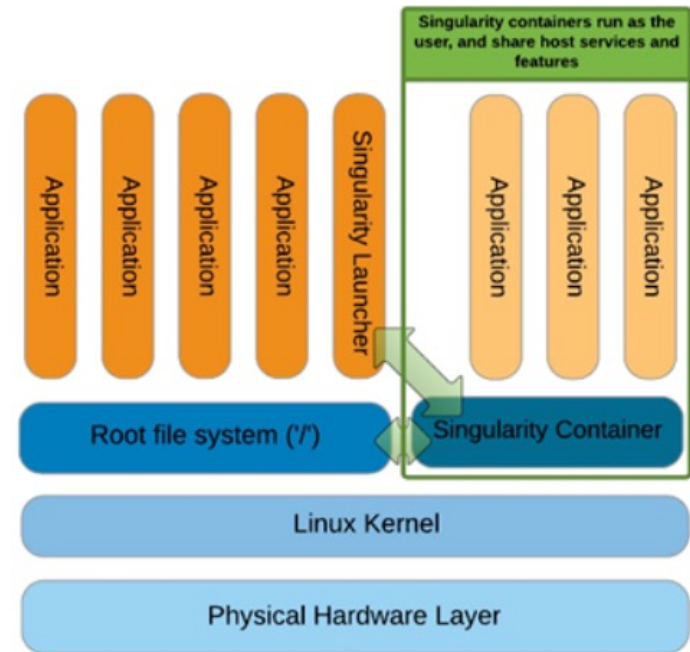


Singularity/Apptainer



Singularity Containers

- No dependency of a daemon
- Can be run as a simple user
 - Avoid permission headaches and hacks
- More easily portable
- Image/container is a file (or directory)



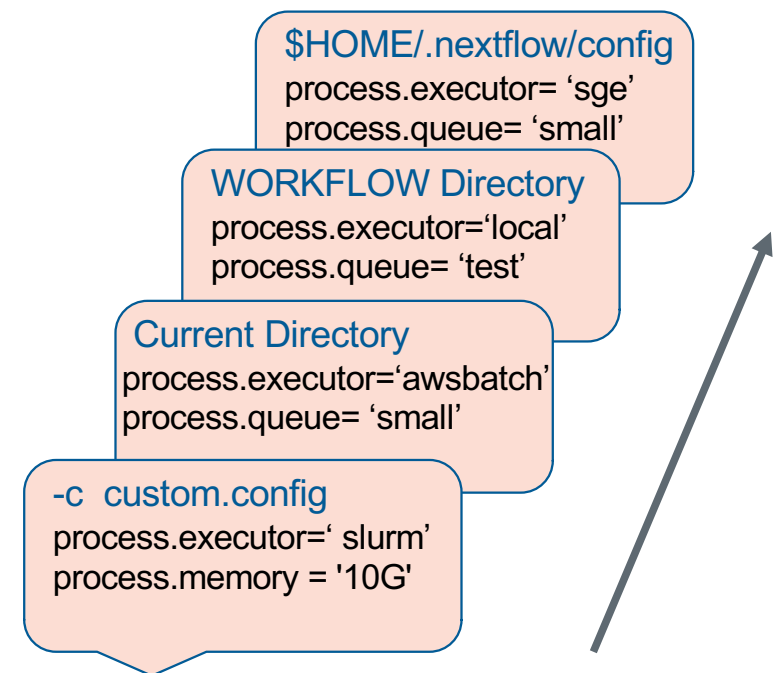
HPC Container
Singularity

Configuring Singularity with Nextflow

Nextflow Configuration File(s)



- Nextflow can load pipeline configurations from multiple locations:
 - Home directory
 - `$HOME/.nextflow/config`
 - workflow project directory
 - `~/.nextflow/assets/nextflow-io/rnaseq/nextflow.config`
 - current directory
 - `$PWD/nextflow.config`
 - config file is given with `-c <config file>`
- Understand the overriding behaviour
 - `process.executor='slurm'`
 - `process.queue= 'small'`
 - `process.memory='10G'`



Configuration Files: scopes



- Configuration settings can be organized in different scopes
- Nextflow scopes
 - *env*
 - *params*
 - *process*
 - *and many other*

```
#scope by dot prefixing
process.executor = 'slurm'
process.queue = 'small'
process.memory = '10G'
```

```
#scope using the curly brackets
singularity {
    enabled = true
    autoMount = true
}
```

Configuration Files: profiles



- A profile is a set of configuration attributes that can be activated when launching a pipeline execution
- Configuration files can contain the definition of one or more profiles.
- Use *-profile* flag to activate attributes *via* command line

```
profiles {  
  
  standard {  
    process.executor = 'local'  
  }  
  
  cluster {  
    process.executor = 'slurm'  
    process.queue = 'small'  
    process.memory = '10.GB'  
  }  
}
```

Configuring Singularity with Nextflow



- Command line interface option : `-with-singularity`
- In 'nextflow.config' file as a *profile* option:

```
singularity {  
  process.container = 'quay.io/nextflow/rnaseq-nf:v1.1'  
  singularity.enabled = true  
  singularity.autoMounts = true  
}
```

Command: `nextflow run main.nf -profile singularity`

Running Nextflow workflows on Puhti

Puhti Recipe for Running Nextflow Pipeline

- ✓ Prepare your singularity images if needed
- ✓ Load Nextflow environment on Puhti
- ✓ Set-up your Nextflow pipeline dependencies
- ✓ Prepare batch job for Nextflow pipeline



Preparing Singularity Images if Needed



- Pull a Singularity image from a singularity registry
 - Use Puhti (e.g., `singularity pull shub://vsoch/hello-world`)
- Convert a Docker image to Singularity one
 - Puhti can work (e.g., `singularity pull docker://tensorflow/tensorflow:latest`)
- Build a Singularity image from scratch
 - Puhti can't be used



Load Nextflow Environment on Puhti



- Puhti uses module system (lmod) to manage software stack
- Nextflow is installed a module
 - `module load nextflow`
- Own installation :
 - `wget -qO- https://get.nextflow.io | bash && mv nextflow ~/bin/`
- One temporarily switch to specific version :
 - `NXF_VER=20.04.0 nextflow run hello ...`

Prepare Your Application Dependencies



- Databases
- Move Singularity images to correct path
- Actual files/samples



Run Nextflow as a Batch Job



- Submit nextflow pipeline as a batch job
- Run all the computations in allocated job
- Avoid using 'slurm' executor from nextflow
- For high-throughput jobs use 'Hyperqueue' executor

```
#!/bin/bash
#SBATCH --job-name=demo_test
#SBATCH --account=project_xxx
#SBATCH --time=48:00:00
#SBATCH --mem-per-cpu=4G
#SBATCH --cpus-per-task=20
#SBATCH --partition=small

# Load a specific version of nextflow
module load nextflow/22.04.5

# clone a specific version of your application and stick to it
nextflow run main.nf -c demo.conf -c cluster.conf \
    -with-singularity ./containers/metaphage.simg -resume

# nextflow run nf-core/sarek -r 2.7.1 \
#     -profile test,singularity -resume
```

Reporting and Visualisation of workflows

Reporting and Visualisation of pipeline



Useful optional flags for creating reports and visualisation

```
-with-dag  
-with-timeline  
-with-report
```

- Execution report

```
nextflow run <nextflow_script> -with-report <file-name>.html
```

- DAG visualisation

```
nextflow run <nextflow_script> -with-dag <file-name>.dot
```

- Timeline report:

```
nextflow run <nextflow_script> -with-timeline <file-name>.html
```

Working with *nf-core* workflows





A community effort to collect a curated set
of analysis pipelines built using Nextflow.

<https://nf-co.re>

nf-core Workflows

- A nice resource for reproducible bioinformatics pipelines
- Provides common pipeline structure and usage
 - Current status: released (49); development (19)
- Each pipeline has its own documentation
 - e.g., `nextflow run nf-core/rnaseq -r 3.0 --help`
- Join on slack/twitter for help



How to Use *nf-core* Workflows ?



- One can use *nextflow* command to fetch a nf-core repository during runtime
 - No need for explicit cloning a repository in advance
 - Nextflow clones nf-core repository to \$HOME directory (~/.nextflow/cache/assets/nf-core...)
- Syntax:

```
nextflow pull nf-core/<pipeline> -r <revision>  
nextflow run nf-core/<pipeline> -r <revision>
```


Running nf-core Workflows at CSC (1/2)



- You don't need to install any nf-core tools to run nf-core pipelines at CSC
 - Just loading nextflow module is enough
- Use singularity as container engines
 - Conda : csc discourages this approach
 - Docker: no root access for users
- Use test profile to see if pipeline works



Running nf-core Workflows at CSC (2/2)



- Change resources (e.g., CPUs, Memory) as needed in production runs
- Containers building can fail in initial attempts
- Explore more by cloning pipeline locally

```
#!/bin/bash
#SBATCH --time=01:00:00
#SBATCH --partition=small
#SBATCH --account=project_xxxx
#SBATCH --cpus-per-task=4
#SBATCH --mem-per-cpu=4000

export SINGULARITYENV_TMPDIR=$PWD
export SINGULARITYENV_CACHEDIR=$PWD
unset XDG_RUNTIME_DIR

# Activate Nextflow on Puhti
nextflow/22.10.1

# nf-core pipeline examples here
# Variant calling on genome data
nextflow run nf-core/sarek -r 2.7.1 -profile test,singularity -resume
# proteomics example
# nextflow run nf-core/proteomicslfq -r 1.0.0 -profile test,singularity -resume
# metabolomics example
# nextflow run nf-core/metaboigniter -r 1.0.1 -profile test,singularity -resume
```

An example batch script for running nf-core pipeline

Debugging nf-core Workflows



- Set Singularity/Apptainer cache/tmp directory explicitly to avoid disk space problems
 - By default, singularity/apptainer cache is in \$HOME directory
- Copying singularity images from galaxyproject can sometimes cause issues
 - Try pulling image from the same host system where you are deploying nf-core pipelines



Debugging nf-core Workflows



- Set \$TMPDIR explicitly to a folder that is writable as sometimes tmpdir
 - Nextflow may not mount it; you need to mount
- Some software have hardcoded /tmp paths; these can't be reset by env variable. Use the following trick
 - --bind writable dir :/tmp



Running Workflows at Scale



- Avoid unnecessary reads and writes of data on Lustre file system to improve I/O performance
 - If unavoidable, use fast local NVMe disk, not Lustre (*i.e.* /scratch)
- Don't run too many/short *job steps* – they will bloat Slurm accounting DB
 - Avoid *slurm* scheduler
- Don't run too long jobs without a restarting option. Increased risk of something going wrong, resulting in lost time/results
 - Use `-resume` flag

.



Running Workflows at Scale



- Don't use Conda installations on Lustre (/projappl, /scratch, \$HOME)
 - Containerize Conda environments instead to improve performance
- Don't create a lot of files, especially within a single folder
 - If you're creating 10 000+ files, you should probably rethink your workflow
 - Consider removing temporary files after job is finished
- Whenever possible, separate serial jobs from parallel ones for efficient usage of resources.

